MTRC

MODERN INSTITUTE OF TECHNOLOGY AND RESEARCH CENTRE

DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING

Faculty Name	Abhimanyu Saraswat		
Subject Name with Code	7CS5-13 Generative AI		
Semester	7th	**Session**	2025-26 (Odd)
Branch	CSE	**Batch**	A,B,AIDS,AIML
Unit No.	4	**Submission Date**	1-Oct-25

Lecture No.	*Topic Covered*	*Page No.*	*Total Pages*
26	Generative models for multimodal data	1 - 2	2
27	Generative models for sequential data	3 - 4	2
28	Advanced techniques:Style transfer	5 - 6	2
29	Advanced techniques:CycleGAN	7 - 8	2

References:

1	Generative Deep Learning: Teaching Machines to Paint, Write, Compose, and Play by David Foster, O'Reilly Media
2	Deep Learning by Ian Goodfellow, Yoshua Bengio, and Aaron Courville
3	Generative Adversarial Networks by Ian Goodfellow, Yoshua Bengio, and Aaron Courville

Faculty HOD Dean Academics

1/10/25 1/10/25

Unit-4 Advanced Generative AI topics

Lecture1

Generative models for multimodal data

Generative models for multimodal data are artificial intelligence systems that can process, generate, and integrate information across various modalities—including text, images, audio, and video—within a unified framework.

Concept and Architecture

Generative multimodal models utilize specific neural network architectures tailored to each data type, such as transformers for text, convolutional neural networks for images, and recurrent models or CNNs for audio and video analysis. These models encode each type of input into a common representation space using dedicated encoders, then combine these representations through fusion layers or attention mechanisms. This integration allows the model to generate outputs that meaningfully combine information from multiple modalities, like describing an image with text, generating an image from a text prompt, or summarizing video content as text.

Capabilities and Applications

Multimodal generative models can:

- Generate images from text descriptions
- Create text captions for images and videos
- Synthesize audio or video sequences based on textual or visual prompts
- Enable cross-modal retrieval, where, for example, an image is found using a text query

Applications include virtual assistants, creative content generation, medical image analysis, adaptive user interfaces, and accessibility features (such as automatic captioning or image descriptions for the visually impaired).

Key Components

- **Encoders:** Transform each modality's input into a vector representation
- **Fusion Mechanisms:** Integrate encoded features from different modalities
- **Decoders:** Generate outputs in the desired modality or modalities
- **Training Datasets:** Require large sets of paired multimodal examples (e.g., image-caption pairs)
- **Loss Functions:** Optimize for cross-modal consistency and output quality

Challenges

Building and training these models is complex, requiring extensive multimodal annotated datasets and computational resources. The fusion of different data types is also non-trivial, as each modality has unique characteristics.

Notable Models

Examples include CLIP (connecting text to images), DALL-E (image generation from text), Stable Diffusion (text-to-image), and Google's AudioPaLM (combining speech and text).

Multimodal generative AI represents a significant advance in AI, enabling more human-like, creative, and context-aware applications across industries.

References

1. Generative Deep Learning: Teaching Machines to Paint, Write, Compose, and Play by David Foster, O'Reilly Media.
2. Deep Learning by Ian Goodfellow, Yoshua Bengio, and Aaron Courville.
3. Generative Adversarial Networks by Ian Goodfellow, Yoshua Bengio, and Aaron Courville.

Lecture2

Generative models for sequential data—such as time series and video—use AI techniques to generate new, realistic data that preserves the temporal or sequential patterns found in the original datasets.

Overview and Key Principles

Generative models for sequential data must account for dependencies across time or sequence steps, unlike standard models for static data. Their core challenge is to learn both the distribution of observations at each time step and the intricate patterns connecting those steps, such as trends, cycles, or causal relationships. This applies across domains like financial forecasting, traffic prediction, physiological monitoring, and video synthesis.

Major Model Types

- **Autoregressive Models:** These generate new sequences step by step, conditioning each new point on previous points. Architectures include Recurrent Neural Networks (RNNs), Long Short-Term Memory (LSTM), and Temporal Convolutional Networks (TCN). Examples include sequence-to-sequence models for video captioning or text generation.

- **Generative Adversarial Networks (GANs):** Time-series GANs (such as TimeGAN and DoppelGANger) have been developed to generate realistic synthetic sequential data by learning both stepwise relationships and global trends. TimeGAN, for instance, incorporates an adversarial loss (like standard GANs) and an explicit supervised loss to ensure the temporal dynamics are accurately captured in the synthetic data.

- **Transformer-Based Models:** Transformers have recently been extended for time series and video, using attention mechanisms to model long-range dependencies in sequential data. These models excel at handling variable-length sequences efficiently and are used in advanced applications like video generation or few-shot time series learning.

Applications

- Synthetic data generation for privacy-preserving analytics
- Video frame synthesis or video prediction
- Financial and energy demand forecasting
- Medical signal generation and augmentation

Notable Example: TimeGAN

TimeGAN is a state-of-the-art model that combines adversarial training with a learned embedding space. It preserves both temporal transitions (conditional probabilities of future

values) and overall data distribution. This model and its variants outperform traditional GANs on measures of realism and predictive utility for time series data.

Advancements

Recently, foundation models for sequential data (e.g., TimesFM) have shown the ability to forecast, generate, or adapt to new patterns with few-shot learning—greatly increasing flexibility in generating or predicting time-dependent data.

Generative models for sequential data thus play a crucial role in advancing synthetic data, prediction, and creative applications spanning time series, audio, and video.

References

1. Generative Deep Learning: Teaching Machines to Paint, Write, Compose, and Play by David Foster, O'Reilly Media.
2. Deep Learning by Ian Goodfellow, Yoshua Bengio, and Aaron Courville.
3. Generative Adversarial Networks by Ian Goodfellow, Yoshua Bengio, and Aaron Courville.

Lecture 3

Advanced style transfer techniques use generative AI to blend the style of one sample (such as an artwork's brushstrokes or a writer's tone) with the content of another image, video, or text. Recent innovations have expanded beyond classic neural style transfer to employ GANs, VAEs, diffusion models, and multimodal methods for higher quality, speed, and cross-domain flexibility.

Major Style Transfer Techniques

- **Neural Style Transfer (NST):** Initiated with convolutional neural networks (CNNs), NST matches the semantic content of one sample (e.g., a photograph) with the style statistics of another (e.g., a painting) using Gram matrix statistics. Adaptive Instance Normalization (AdaIN) improved speed and allowed real-time style manipulation.

- **GAN-Based Style Transfer:** Generative Adversarial Networks (GANs) made style transfer more photorealistic by having a generator produce stylized outputs and a discriminator assess their realism. GANs introduced latent space modeling and disentanglement techniques, leading to greater output diversity and control over the separation of content and style.

- **Variational Autoencoders (VAEs):** VAEs facilitate structured latent space learning, providing control and interpolation between styles. While not always as photorealistic as GANs, they are efficient at preserving global structure and enabling smooth transitions between styles.

- **Diffusion-Based Style Transfer:** Diffusion and autoregressive models have become state-of-the-art for high-resolution, fine-detail style transfer. They iteratively add and remove noise, enabling flexible conditioning on text or other attributes, and they achieve exceptional sample quality and stability compared to GANs, though with higher computational demands.

- **Multimodal & Text-Based Style Transfer:** Techniques now allow transferring style between different data types, such as images to 3D, or using text prompts to guide style adaptation (as seen with models leveraging CLIP for semantic control).

Applications

- **Images and Video:** High-resolution stylization, real-time video style transfer, portrait retouching, artistic effects, and animation.

- **Text:** Adapting tone, sentiment, or writing style in natural language processing for creative generation and targeted marketing.

- **Other Domains:** 3D asset creation, game design, and multimodal art synthesis.

Ongoing Challenges

Despite progress, tradeoffs remain between computational efficiency, content fidelity, and controllable diversity. Modern systems emphasize evaluation metrics such as perceptual realism, semantic preservation, and style diversity to assess effectiveness.

Advanced style transfer thus leverages a variety of generative architectures—each with strengths and constraints—to create increasingly powerful, flexible, and controllable transformations across visual, textual, and cross-modal domains.

References

1. Generative Deep Learning: Teaching Machines to Paint, Write, Compose, and Play by David Foster, O'Reilly Media.
2. Deep Learning by Ian Goodfellow, Yoshua Bengio, and Aaron Courville.
3. Generative Adversarial Networks by Ian Goodfellow, Yoshua Bengio, and Aaron Courville.

Lecture 4

CycleGAN (Cycle-Consistent Generative Adversarial Network) is an advanced technique for unpaired image-to-image translation, enabling style or domain transfer between two sets of images without requiring paired training data.

Core Architecture and Mechanisms

- CycleGAN employs two generators, G:X→YG:X→Y and F:Y→XF:Y→X, where XX and YY are image domains. GG converts images from domain XX to YY, while FF attempts to invert that transformation.

- Two discriminators, DXDX and DYDY, independently learn to distinguish real images from generated images in each domain.

- The key innovation is the **cycle consistency loss**, which ensures that applying GG followed by FF (or vice versa) reconstructs the original image, enforcing meaningful bijective mappings between domains:

$x \approx F(G(x)), y \approx G(F(y))x \approx F(G(x)), y \approx G(F(y))$

- The generators have encoder-transformer-decoder architectures using convolutional layers and residual blocks. PatchGAN discriminators evaluate image patches instead of entire images, focusing on local realism.

Advanced Techniques and Improvements

- **Attention Mechanisms:** Integrating attention modules into CycleGAN enhances the model's ability to focus on significant regions, improving translation quality and detail preservation, especially in face image generation and style transfer tasks.

- **Normalization Techniques:** Adaptive normalization strategies like AdaLIN help the network handle shape and texture changes more flexibly during image translation.

- **Single-Domain Generalization:** Recent variants (e.g., SDCGAN) extend CycleGAN to fault diagnosis and other applications by enabling generalization from a single domain to multiple unseen domains.

- **High-Resolution and Efficient Training:** Techniques such as multi-scale architectures and optimized training schedules improve output resolution and training stability.

- **Medical and Low-Light Imaging:** CycleGAN adaptations enhance image quality for domain-specific applications like MRI-to-CT translation and low-light video enhancement, demonstrating CycleGAN's flexibility across domains.

Applications

- Style and seasonal changes in photos
- Converting artwork styles or sketches to colored images

- Medical image translation and augmentation
- Video enhancement and low-light correction
- Fault diagnosis and anomaly detection in engineering domains

CycleGAN remains a powerful framework because of its flexibility in unsupervised domain translation and the wide array of improvements through architectural and application-specific innovations.CycleGAN (Cycle-Consistent Generative Adversarial Network) is an advanced generative modeling technique designed for unpaired image-to-image translation between two domains without requiring paired data. It works with two generators (one mapping domain X to Y and another mapping back from Y to X) and two discriminators that ensure generated images look realistic in each domain. The core innovation is the cycle consistency loss which enforces that mapping an image from one domain to the other and back should reconstruct the original image, thus preserving content while changing style or domain representations.

Advanced CycleGAN techniques include:

- Using attention mechanisms to improve focus on important features for image generation.
- Adaptive normalization methods like AdaLIN that help handle texture and shape changes flexibly.
- Single-domain generalization extensions to allow domain adaptation from one source domain to new unseen domains.
- Multi-scale architectures and optimized training helping to produce high-resolution outputs and improve stability.
- Specialized versions for medical imaging, low-light enhancement, and fault diagnosis in engineering applications.

These advancements make CycleGAN highly effective for tasks such as style transfer, seasonal changes in photos, medical image modality translation, sketch to color, and video enhancement, proving flexible and robust across varied computer vision challenges.

References

1. Generative Deep Learning: Teaching Machines to Paint, Write, Compose, and Play by David Foster, O'Reilly Media.
2. Deep Learning by Ian Goodfellow, Yoshua Bengio, and Aaron Courville.
3. Generative Adversarial Networks by Ian Goodfellow, Yoshua Bengio, and Aaron Courville.